



**FACULTY
OF MATHEMATICS
AND PHYSICS**
Charles University

BACHELOR THESIS

Michal Pospěch

Thesis title is N/A

Department of Theoretical Computer Science and Mathematical Logic

Supervisor of the bachelor thesis: **Mgr. Roman Neruda, CSc.**

Study programme: **Computer science**

Study branch: **General computer science**

Prague **2021**

I declare that I carried out this bachelor thesis independently, and only with the cited sources, literature and other professional sources. It has not been used to obtain another or the same degree.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In date

Author's signature

Dedication. It is nice to say thanks to supervisors, friends, family, book authors and food providers.

Title: Thesis title is N/A

Author: Michal Pospěch

Department: Department of Theoretical Computer Science and Mathematical Logic

Supervisor: Mgr. Roman Neruda, CSc., Department of Theoretical Computer Science and Mathematical Logic

Abstract: Abstracts are an abstract form of art. Use the most precise, shortest sentences that state what problem the thesis addresses, how it is approached, pinpoint the exact result achieved, and describe the applications and significance of the results. Highlight anything novel that was discovered or improved by the thesis. Maximum length is 200 words, but try to fit into 120. Abstracts are often used for deciding if a reviewer will be suitable for the thesis; a well-written abstract thus increases the probability of getting a reviewer who will like the thesis.

Keywords: key words

Contents

Introduction	3
1 Important first chapter	5
2 Theory	7
2.1 Reinforcement learning	7
2.2 Evolutionary algorithms	7
2.3 Evolutionary strategies	8
2.3.1 CMA-ES	9
2.4 Evolutionary strategies as replacement for reinforcement learning	9
2.4.1 OpenAI Evolutionary Strategy	10
2.4.2 Novelty search	11
3 Results and discussion	15
Conclusion	17
Bibliography	19
A Using CoolThesisSoftware	21

Introduction

Chapter 1

Important first chapter

Chapter 2

Theory

2.1 Reinforcement learning

- problem description (state space, action space, reward...)
- various methods
 - value function
 - criterion of optimality
 - direct policy search (and various methods)

2.2 Evolutionary algorithms

Evolutionary algorithms (EA) are a type of optimisation metaheuristics inspired by the process of biological evolution. At first a number of possible solutions to the problem at hand is generated (*population*) and each solution (*individual*) is encoded (via a domain-specific encoding) and evaluated giving us the value of its *fitness*. Fitness is a function describing how good that particular individual is and it is everything that is needed for creation of Then a new population is created using a *crossover* (combination) of 1 or more individuals which are selected using the operator of *parental selection*. Each of the newly created individuals has a chance to be mutated via the *mutation* operator. Finally a new population is selected from *offsprings* and possibly the parents based of fitness and enters the next iteration of the EA and the following generation is chosen using *environmental selection* operator. The algorithm repeats until the stop condition is met, usually a set number of iterations or small improvement of fitness between 2 generations.

There are many variants of EAs such as genetic algorithms (most common), genetic programming, evolutionary programming, neuroevolution or evolutionary strategies that are further described in following chapter. [1] [2]

Algorithm 1 Evolutionary algorithm

```

1: initialize population  $P^0$  with  $n$  individuals
2: set  $t = 0$ 
3: repeat
4:    $Q^t = \{\}$ 
5:   for  $i \in \{1 \dots m\}$  do
6:      $p_1, \dots, p_\rho = \text{ParentalSelection}(P^t)$ 
7:      $q = \text{Crossover}(p_1, \dots, p_\rho)$ 
8:      $q = \text{Mutation}(q)$  with chance  $p$ 
9:      $Q^t = q \cup Q^t$ 
10:  end for
11:   $P^{t+1} = \text{EnvironmentalSelection}(Q^t \cup P^t)$ 
12:  increment  $t$ 
13: until stop criterion fulfilled

```

2.3 Evolutionary strategies

Evolutionary strategies (ES) are a type of optimisation metaheuristic which further specialises EA and restricts their level of freedom. The selection for crossover is unbiased, mutation is parametrised and thus controllable, individuals which should be put to next generation are chosen ordinaly based on fitness and individuals contain not only the problem solution but also control parameters.

More formally ES $(\mu/\rho, \kappa, \lambda)$ has μ individuals in each generation, which produces λ offsprings, each created by crossover of ρ individuals and each individual is able to survive for up to κ generations as described in algorithm 2. This notation further generalizes the old (μ, λ) and $(\mu + \lambda)$ notations, where the " " notation means $\kappa = 1$ and "+" notation $\kappa = \infty$.

To design an ES one must first select an appropriate representation for an individual and the most natural one is preferred in most cases, if all parameters are of one type (e.g. a real number) a simple vector will suffice, if the types are mixed, a tuple of vectors is required. This however causes an increased complexity of the variation operator.

As for design of the variation operator there are some guidelines that should be followed when designing it.

Algorithm 2 $(\mu/\rho, \kappa, \lambda)$ -ES

```
1: initialize population  $P^0$  with  $\mu$  individuals
2: set age for each  $p \in P^0$  to 1
3: set  $t = 0$ 
4: repeat
5:    $Q^t = \{\}$ 
6:   for  $i \in \{1 \dots \lambda\}$  do
7:     select  $\rho$  parents  $p_1, \dots, p_\rho \in P^t$  uniformly at random
8:      $q = \text{variation}(p_1, \dots, p_\rho)$  with age 0
9:      $Q^t = q \cup Q^t$ 
10:  end for
11:   $P^{t+1} = \text{select } \mu \text{ best (wrt. fitness) individuals from } Q^t \cup \{p \in P^t : \text{age}(p) < \kappa\}$ 
12:  increment age by 1 for each  $p \in P^{t+1}$ 
13:  increment  $t$ 
14: until stop criterion fulfilled
```

Reachability every solution should be reachable from any other solution in a finite number of applications of the variation operator with probability $p > 0$

Unbiasedness the operator should not favour any particular subset of solution unless provided with information about problem at hand

Control the operator should be parametrised in such way that the size of the distribution can be controlled (practice had shown that decreasing it as the optimal solution is being approached is necessary)

[3] [1]

kovariance

2.3.1 CMA-ES

TODO [4]

2.4 Evolutionary strategies as replacement for reinforcement learning

Black-box optimisation is an alternative approach to solving RL tasks also known as Direct policy search or neurevolution when applied to neural networks. It has

several attractive properties such as indifference to distribution of rewards, no need for backpropagation and tolerance of arbitrarily long episodes.

Compared to reinforcement learning using evolutionary strategies has the advantage of not needing a gradient of the policy performance. Also as the state transition function is not known the gradient can't be computed using backpropagation-like algorithm. Thus some noise needs to be added to make the problem smooth and the gradient to be estimable. Here is where reinforcement learning and evolutionary strategies differ, reinforcement learning adds noise in the action space (actions are chosen from a distribution) while evolutionary strategies add noise in the parameter space (parameters perturbed while actions are deterministic).

Not requiring backpropagation has several advantages over other RL methods. First the amount of computation necessary for one episode of ES is much lower (about one third, potentially even less for memory usage). Not calculating gradient using analytic methods also protects these methods from suffering from *exploding gradient* which is a common issue with recurrent neural networks. And last, the network can contain elements that are not differentiable such as hard attention.

Also ES are easily parallelisable. First, contrary to RL, where value function is inherently linear procedure and has to be performed more times to improve a given policy. Furthermore it operates on whole whole episodes, therefore it does not require frequent communication between workers. And finally, as the only information received by each worker is the reward, it is possible to communicate only the reward inbetween workers. That, however, requires synchronising seeds known to other workers beforehand and recreating perturbations based on that information. Thus the required is extremely low compared to communication of whole gradients which would be required for parallelisation of a policy gradient based algorithm.

2.4.1 OpenAI Evolutionary Strategy

These ideas are explored in OpenAI-ES algorithm. The agent acting in the environment is represented by a policy π_θ with parameter vector θ and $f(\cdot)$ is the reward returned by the environment. As we need to introduce some noise, the population distribution p_{ψ} is instantiated as an isotropic multivariate Gaussian with mean ψ and covariance $\sigma^2 I$. Then $\mathbb{E}_{\theta \sim p_\psi} f(\theta) = \mathbb{E}_{\epsilon \sim N(0, I)} f(\theta + \sigma \epsilon)$. Thus the gradient approximation is calculated as follows:

$$\nabla_{\theta} \mathbb{E}_{\theta \sim p_\psi} f(\theta) = \nabla_{\theta} \mathbb{E}_{\epsilon \sim N(0, I)} f(\theta + \sigma \epsilon) \approx \frac{1}{n\sigma} \sum_{i=1}^n f(\theta_i) \epsilon_i, \quad (2.1)$$

where $\theta_i = \theta + \sigma \epsilon_i$, $\epsilon_i \sim N(0, I)$ and n is the population size.

The resulting algorithm uses SGD (or Stochastic Gradient Ascent - SGA in this case) or other gradient based optimisation technique for parameter vector θ update.

As the ES could be seen as method for computing a derivative estimate using finite differences in randomly chosen direction it would suggest that it would scale poorly with dimensions of parameters θ same as the finite differences method. In theory the number of necessary optimisation steps should scale linearly with the dimension. That however doesn't mean that larger networks optimised using ES will perform worse than smaller ones, that depends on the difficulty (intrinsic dimension) of the problem. The network will perform the same however it will take more optimisation steps to do so.

In practice ES performs slightly better on larger networks and it is hypothesised that it is for the same reason as why it is easier to optimise large networks using standard gradient based methods: larger networks have fewer local minima.

Due to perturbing the parameters and not the actions ES are invariant to the frequency at which the agent acts in the environment. Traditional MDP-based reinforcement learning methods rely on *frameskip* as one their parameters that is crucial to get right for the optimization to be successful. While this is solvable for problems that do not require long term planning and actions, long term strategic behaviour poses a challenge and reinforcement learning needs hierarchy to be succesful unlike evolutionary strategy.

2.4.2 Novelty search

While the main drive in of improvement in ES is the value of fitness (how "good" the result is), novelty search takes a different approach. Novelty search is focused on finding different solutions, as it is inspired by nature's drive towards diversity. Each policy has its novelty calculated with respect to previous policies and search is directed to parts of search space with high novelty. This approach makes it less susceptible to local optima created by deceptive rewards than reward-based method.

Each policy π gets assigned its domain-dependent behavioral characteristics $b(\pi)$ (e.g. final position of the agent) and it is added to an archive set A of characteristics of previous policies. Then the novelty $N(b(\pi_\theta), A)$ is calculated as average distance from k nearest neighbours from the archive set A .

$$N(\theta, A) = N(b(\pi_\theta), A) = \frac{1}{|S|} \sum_{j \in S} \|(\pi_\theta) - b(\pi_j)\|_2$$

$$S = kNN(b(\pi_\theta), A)$$

Combination with evolution strategies

To find and follow the gradient of expected novelty with respect to θ^t we use the framework outlined in 2.4.1.

With archive A and sampled parameters $\theta_t^i = \theta_t + \sigma\epsilon_i$, the gradient estimate can be calculated via following formula:

$$\nabla_{\theta_t} \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} [N(\theta_t + \sigma\epsilon) | A] \approx \frac{1}{n\sigma} \sum_{i=1}^n N(\theta_t^i, A) \epsilon_i \quad (2.2)$$

It is possible because archive A is fixed during one iteration and is updated only at the end. Only characteristics corresponding to each θ^t are added to A , as adding each sampled would cause the archive A to inflate too much increasing the complexity of calculation of nearest-neighbours.

To encourage additional diversity an initial meta-population of M , selection of M is domain dependent, agents is created. While it is possible to optimise the behaviour of a single agent and reward it for behaving differently than its ancestors, this way we get the benefits of population-based exploration. Each agent has parameters θ^m and is being rewarded for behaviour different from all prior agents, thus we get M differently behaving policies.

M random parameter vectors are initialised and in each iteration one is selected to be updated. The selection probability is proportional to its novelty.

$$P(\theta^m) = \frac{N(\theta^m, A)}{\sum_{i=1}^M N(\theta^i, A)} \quad (2.3)$$

To perform the update step, we need to calculate the gradient estimate of expected novelty with respect to θ_t^m using equation 2.2 where n is the number of perturbations. When we get the gradient estimate we use SGD (or in this case Stochastic Gradient Ascent) with learning rate α to update the parameters θ^m

$$\theta_{t+1}^m := \theta^m + \alpha \frac{1}{n\sigma} \sum_{i=1}^n N(\theta_t^{i,m}, A) \epsilon_i. \quad (2.4)$$

After updating the individual, a new behavioral characteristics $b(\pi_{\theta_{t+1}^m})$ is calculated and added to the archive A .

This process is repeated for a predetermined number of times as novelty search is not supposed to converge to a "best" solution and returns the best performing policy which is being preserved during the run of the algorithm.

Combination with reward based exploration

While *NS-ES* helps agents avoid local optima and deceptive reward signals, it also completely discards reward which might cause the performance to suffer.

Describe?

Therefore NSR-ES, an improved version of NS-ES, uses both reward (fitness) and novelty for computation of the update step. NSR-ES is in many ways similar to NS-ES, it calculates both novelty and reward at once and it operates on entire episodes. The only difference is, that the calculation of the gradient estimate is based on the average of reward and novelty.

Specifically, for parameter vector $\theta_t^{m,i} = \theta_t^m + \sigma \epsilon_i$ we calculate the reward $f(\theta_t^{m,i})$ and novelty $N(\theta_t^{m,i}, A)$, rank-normalise both values independently (as both values usually have completely different scales), calculate the average and set it as weight for corresponding ϵ_i for gradient estimation. Then the estimated gradient is used to update the parameter vector via SGD (or other gradient based optimisation method) similarly as in equation 2.4:

$$\theta_{t+1}^m := \theta_t^m + \alpha \frac{1}{n\sigma} \sum_{i=1}^n \frac{N(\theta_t^{i,m}, A) + f(\theta_t^{i,m})}{2} \epsilon_i. \quad (2.5)$$

Intuitively, following the approximated gradient based on both novelty and reward directs the search areas of the parameter-space with both high novelty and reward. This can, however, be improved further.

While NSR-ES uses a linear combination of reward and novelty to approximate that is static for the whole duration of the training. Contrary to that NSRAdapt-ES (NSRA-ES) dynamically changes the ratio of reward gradient $f(\theta_t^{i,m})$ and novelty gradient $N(\theta_t^{i,m}, A)$ based on how the training is currently progressing. This way, it will give more weight to the reward (thus following the performance gradient) when making progress and more weight to the novelty (following novelty gradient) when stuck in a local optima to give more incentive to find different approaches.

Formally, a parameter w is used to control the ratio of reward and novelty used for calculation of gradient estimate. For a specific w at a given generation parameter vector θ_t^m is updated (SGD used) via following expression:

$$\theta_{t+1}^m := \theta_t^m + \alpha \frac{1}{n\sigma} \sum_{i=1}^n (1-w)N(\theta_t^{i,m}, A)\epsilon_i + wf(\theta_t^{i,m})\epsilon_i. \quad (2.6)$$

[5]

- Evolution Strategies as a Scalable Alternative to Reinforcement Learning

[6]

- algorithm description
- comparison with RL
- parallelization

- Improving Exploration in Evolution Strategies for Deep Reinforcement Learning via a Population of Novelty-Seeking Agents [5]
 - ratio of fitness and novelty and its effects

Chapter 3

Results and discussion

Conclusion

Bibliography

- [1] Günter Rudolph. “Evolutionary Strategies”. In: *Handbook of Natural Computing*. Ed. by Grzegorz Rozenberg, Thomas Bäck, and Joost N. Kok. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 673–698. ISBN: 978-3-540-92910-9. doi: 10.1007/978-3-540-92910-9_22. URL: https://doi.org/10.1007/978-3-540-92910-9_22.
- [2] P. A. Vikhar. “Evolutionary algorithms: A critical review and its future prospects”. In: *2016 International Conference on Global Trends in Signal Processing, Information Computing and Communication (ICGTSPICC)*. 2016, pp. 261–265. doi: 10.1109/ICGTSPICC.2016.7955308.
- [3] Hans-Paul Schwefel and Günter Rudolph. “Contemporary evolution strategies”. In: *Advances in Artificial Life*. Ed. by Federico Morán et al. Berlin, Heidelberg: Springer Berlin Heidelberg, 1995, pp. 891–907. ISBN: 978-3-540-49286-3.
- [4] Nikolaus Hansen. *The CMA Evolution Strategy: A Comparing Review*. 2006.
- [5] Edoardo Conti et al. *Improving Exploration in Evolution Strategies for Deep Reinforcement Learning via a Population of Novelty-Seeking Agents*. 2018. arXiv: 1712.06560 [cs.AI].
- [6] Tim Salimans et al. *Evolution Strategies as a Scalable Alternative to Reinforcement Learning*. 2017. arXiv: 1703.03864 [stat.ML].

Appendix A

Using CoolThesisSoftware

